


90-Day Execution Plan

Outbound machine, tool stack, weekly workflows, and the numbers I'd report.

Saurabh Navale · 24 September 2026 · Companion to the GTM Blueprint and Sales Playbook

Day 1 is assumed to be 1 October 2026. Shift the dates, not the sequence. Every diagram in this document is reproducible in Notion as a Mermaid block; the source markdown carries the code.

1. The Whole Machine



The design decision: the warm path is checked **before** any cold sequence runs. A Happenstance hit converts several times better than the best cold email I can write, and it costs one query.

2. Tool Stack and Budget


Bootstrapped constraint: **no ad spend, no data-vendor contracts.** Everything below is month-to-month and cancellable.

TOOL	JOB IN THE MOTION	COST / MONTH	WHY THIS ONE
DIVYAM.AI - 90-DAY EXECUTION PLAN			SAURABH NAVALE - 24 SEP 2026
Happenstance	Network search — "who do we know at Swiggy in platform engineering?" Warm-path discovery across Gmail, LinkedIn and calendar	\$30 / user	Checked first on every account. Warm beats cold at any volume. SOC 2, API access
SalesRobot	LinkedIn automation for Tier 2 — connection and follow-up sequences, AI icebreakers from recent posts, mini-CRM	\$59–99	Safe-mode throttling and adaptive limits. A restricted account costs more than the subscription
Smartlead	Cold email infrastructure — unlimited inboxes, unlimited warmup, rotation across domains	\$39–94	Unlimited mailboxes on every tier is the whole reason. Warmup included
Sales Navigator	Account and persona mapping, trigger events, hiring signals	~\$99	Non-negotiable for Indian platform-engineering orgs — org charts aren't in any database
Apollo	Contact discovery and enrichment baseline	\$49–79	Start here; assume patchy India coverage and verify manually
Domains + inboxes	2 domains × 3 inboxes. The primary domain never sends cold	~\$15	Reputation firewall
Verification	Bounce control before every send	~\$20	Above 3% bounce, deliverability collapses
CRM	HubSpot free tier or Attio	\$0–29	Stages and required fields matter more than the logo
Total	Against a ₹0 media budget	\$320–465 ₹27k–39k	

Deliberate omission — Clay at \$185+/month is the obvious upgrade for waterfall enrichment. I'd defer it to month 3, and only if manual verification proves to be the bottleneck. At 200 contacts a month, my time is cheaper than the subscription.

3. Channel Workflows

3.1 WARM PATH FIRST — HAPPENSTANCE



Rule: a Happenstance hit never enters an automated sequence. Burning a warm path with a templated message is the most expensive mistake in this plan.

SAURABH NAVALE · 24 SEP 2026

Sandeep's ask, made concrete: I'd run the network query across all 50 accounts in week 1, then bring a **ranked list of ~10 intro requests with the message pre-drafted**. His cost is forwarding an email, not thinking of names.

3.2 LINKEDIN — SALESROBOT, TIER 2

STEP	DAY	ACTION	CEILING
1	D1	Connection request, no note for cold profiles	20–25 / day
2	D3 after accept	Value message — one observation about their AI feature. No ask	—
3	D6	The disqualifying question	—
4	D10	Artifact — MakeMyTrip teardown or a published benchmark	—
5	D14	Soft close, or exit to nurture	—

Safety discipline. Stay at 20–25 connects a day even though the tool allows more. Warm the account for two weeks before automating. Never run SalesRobot and manual outreach on the same profile on the same day. **A restricted profile costs a month of pipeline.**

Where AI personalisation helps, and where it doesn't: SalesRobot's icebreakers drawn from recent posts are fine for step 1. Steps 2 and 3 carry a technical claim and get written by me — an AI icebreaker that misreads an engineering post is worse than no personalisation at all.

3.3 EMAIL — SMARTLEAD



Capacity at steady state is **6 inboxes × 35/day ≈ 200 sends a day** — far beyond what 50 accounts need. The constraint is deliberately list quality, not send volume.

3.4 TIER 1 — NO AUTOMATION AT ALL

Fifteen accounts. Hand-written, one at a time, with a modelled inference-spend estimate in the first email. Roughly 25 minutes each. **These are the accounts that become the case studies.**

4. Weekly Operating Rhythm

DIVYAM.AI · 90-DAY EXECUTION PLAN

SAURABH NAVALE · 24 SEP 2026

FIGURE 4 — THE WEEK, BLOCKED



Daily floor: 25 touches · 5 new accounts researched · inbound answered inside 4 working hours · no deal left in a stage past its clock. **Prospecting hours are meeting-free by policy** — they are the only hours that build next month.

5. The 90 Days


FIGURE 5 — SCHEDULE · DAY 1 = 1 OCTOBER 2026



WEEK BY WEEK

WEEK	FOCUS	90-DAY EXECUTION PLAN	OUTPUT BY FRIDAY	SAURABH NAYAK 21 SEP 2026	LEADING METRIC
1	Stack + list	Tools live, CRM stages configured, domains warming, 50 accounts identified	50 accounts, 0 sends		
2	Data + warm paths	200 contacts verified and tiered; Happenstance map complete; 10 pre-drafted intro requests to Sandeep	200 contacts		
3	Tier-1 wave	15 hand-written emails with spend models; LinkedIn sequences start	First replies		
4	Volume on	Email sequences live post-warmup; all five touches running	≥5 discovery calls		
5	Convert	Technical deep-dives; first pilot scoping docs drafted	2 deep-dives		
6	First pilots	1–2 pilots signed ; cost-audit page live	1 signed pilot		
7	Run + refine	Baselines captured; ICP tightened on evidence — double down on the segment that replies	Pilot baseline set		
8	Second wave	Refreshed list for responding segments; benchmark one drafted	15 total discoveries		
9	Pilot readouts	First ₹-saved number produced; champion prepped to present internally	First verified saving		
10	Business case	Procurement and security started on pilot one; benchmark one published	Security review open		
11	Close motion	Pricing negotiated against a real number; pilots two and three in readout	First close in sight		
12	Close	First logo signed ; case study drafted	1 logo		
13	Systematise	POC runbook written; Q2 pipeline built; three logos, or a clear read on why not	3 logos		

6. Funnel Math, With Capacity



7. Pilot Workflow — Who Does What



The rule that makes this work: the champion presents, not me. A number defended by their own engineer survives a budget meeting. A vendor deck does not.

Success metric wording that works: "≥30% reduction in blended cost per request on [workload], with rubric score within 2% of baseline, measured over 14 days of production traffic."

8. Risk Register

DIVYAM.AI · 90-DAY EXECUTION PLAN

SAURABH NAVALE · 24 SEP 2026

RISK	SIGNAL IT'S HAPPENING	MITIGATION, PRE-PLANNED
LinkedIn restriction	Warning, or invites stop landing	Stay at 20–25/day; warm two weeks; pause automation for seven days
Domain burn	Bounce above 3%, replies fall off a cliff	Secondary domains only; verify before every send; re-warm on trigger
India data quality	Apollo returns stale titles for platform-engineering roles	Sales Navigator as source of truth; Tier 1 verified 100%, Tier 2 sampled
Pilot capacity ceiling	More pilots scoped than engineering can staff	Pace outbound; queue pilots with dates rather than starting them badly
BFSI cycle length	Security review runs past six weeks	Start the security questionnaire at pilot kickoff, not at close
No warm paths	Happenstance returns nothing across the list	Fall back to community-led credibility — benchmarks and OSS, started in parallel anyway
Wrong ICP	Replies cluster in a segment I ranked low	Week 7 is the scheduled pivot point, decided on reply data rather than instinct

9. What I'd Show You at Each Review

REVIEW	THE ARTIFACT	THE QUESTION I'D EXPECT
Day 30	Account map, reply rate by segment and persona, five or more discoveries booked, warm-path results	"Is the message landing, and with whom?"
Day 60	Three pilots running with signed scopes, first baseline numbers, ICP narrowed to two segments	"Are the pilots producing numbers, or activity?"
Day 90	Three logos, one case study, POC runbook, Q2 pipeline, pricing validated in three real negotiations	"Can the next AE run this without you?"

The one number I'd put on the wall: accounts where I can state their monthly inference spend. Everything downstream follows it — it is the difference between a pipeline and a list.

The Six Questions, Still Open

DIVYAM.AI · 90-DAY EXECUTION PLAN

SAURABH NAVALE · 24 SEP 2026

01 · Pricing	Share of verified savings, platform fee, or per-token? It sets ACV and who is in the room.
02 · Capacity	How many concurrent pilots can engineering genuinely support? This plan's ceiling.
03 · Product	Does the router work per-step inside agent workflows, or only per-prompt?
04 · Trust	Security questionnaire pack or SOC 2 path — build, buy, or route around it with air-gap?
05 · Moat	When Portkey ships evidence-based routing, what still can't they copy?
06 · Win	At Day 90 — logos, revenue, or a repeatable motion? Your priority sets mine.