

FOUNDING ACCOUNT EXECUTIVE · INDIA

Go-To-Market Blueprint for Divyam.AI

A 90-day plan to turn a closed-loop routing architecture into signed enterprise revenue — without a marketing budget.

PREPARED FOR

Sandeep Kohli, CEO

PREPARED BY

Saurabh Navale

DATE

22 September 2026

The argument in five lines.

- Divyam is not an observability vendor. Observability is a cost centre. **Divyam is a P&L line item** — and that changes who signs.
- The wedge is the **closed loop**: EvalMate sets the quality bar, Model Router enforces it, traffic moves with **no engineer in the approval path**.
- The proof is already public: **MakeMyTrip cut LLM cost 63% with zero quality loss**. That number opens every first call.
- We will not outspend **Portkey (\$15M Series A)**. We out-architect them, and we out-work them on a narrow, named account list.
- Two engines, unequal weight: **outbound delivers the first three logos; inbound is planted now to compound in Q2**.

90-day commitment: 40 qualified technical discoveries → 12 paid single-workload pilots → 3 closed logos.

Inference cost scales with usage, not headcount.

Every successful AI feature makes gross margin worse. The arbitrage to fix it is sitting unclaimed on the table.

8–20×

Price spread between frontier and open-weight models on the **same workload**

\$2.50 → \$0.30

GPT-4o vs DeepSeek-flash, per 1M input tokens.
Llama 4 Scout: \$0.10

~300×

Token price decline since 2023 — and enterprise bills still went **up**

- Nobody moves traffic to the cheaper model, because **nobody can prove quality holds**. The blocker is evidence, not price.
- Divyam already named the disease: **model inertia**. It is the most commercially useful phrase we own — it turns inaction into a quantifiable monthly cost.

Two markets. Two entirely different buyers.

UNITED STATES

- Heavily funded, converging on **broad AI observability** — traces, spans, scores, dashboards.
- Braintrust raised one of 2026's largest eval rounds; LangSmith owns the LangChain-native install base.
- Buyer: the **AI engineer debugging quality**. Budget: tooling line.

INDIA

- **Hyper cost-conscious**. Buys ROI inside a quarter, not a platform vision inside a year.
- **Sovereignty is contractual**. DPDP plus RBI payment-data localisation gate every BFSI conversation before feature one.
- DPDP programmes alone run **₹15–60 lakh** at large enterprises — we compete for that same budget line.
- Buyer: the **executive defending margin**. Budget: cost of goods.

Consequence for our roadmap: **air-gapped deployment is not a feature. It is the entry ticket to the top 50 accounts in this market.**

Six credible players. One structural gap.

PLAYER	FUNDING & SCALE	WHERE THEY LOSE TO US
Braintrust	Largest 2026 eval-category round	Scores tell you what broke. A human still has to act on them.
LangSmith	LangChain-native, enterprise tier	Locked to LangChain-shaped stacks. Observability-first, never cost-first.
TruEra / Arize	Acquired · enterprise-scale	ML-era monitoring posture. Not built to move live GenAI traffic.
Portkey.ai	\$15M Series A · 500B+ tokens · 125M daily requests · 24k orgs	Routing rules are static and human-authored . Fallbacks are failover, not evidence-based promotion.
Maxim AI	\$3M seed · Elevation Capital	Evals are a destination , not a control signal. No traffic control plane.
Simplismart	\$7M Series A · ~\$20M Nvidia-led round reported	Optimises how a model runs. Never asks which model should run.

Every one of them makes a human the bottleneck between knowing and saving.

Stated plainly, so we don't pretend otherwise.

- We lose **every** paid-search, conference-booth and content-volume fight. Portkey has fifteen million reasons to win those.
- So we don't enter them. **Zero top-of-funnel spend for 90 days.** Every rupee goes to data, tooling and POC engineering.
- We win on **architecture** and on **reps**: 100 surgical, technically literate touches beat 100,000 impressions.
- Three unfair advantages, and only three: **the closed loop, air-gapped deployment, and a POC shipped the same week it is scoped.**
- Being unfunded is also a message: **we have no incentive to sell shelfware.** We get paid when the invoice drops.

Everyone else measures quality.

Only Divyam lets quality **automatically authorise** a cheaper model.

01 · BASELINE

EvalMate

An agentic workflow turns **your** examples into a structured rubric — not a public benchmark. The aligned judge distils into an **~8B reward model running on your own infrastructure**.



02 · ENFORCE

Model Router

A lightweight classifier routes every prompt across **100+ models** against your cost ceiling and quality floor. Trivial work drops to cheap models; hard work escalates to frontier.



03 · CLOSE THE LOOP

Route → Evaluate → Optimise

A new model lands Tuesday. It is benchmarked against your real traffic overnight and **promoted by morning — only if the rubric holds**.

WHAT WE SELL

- **Automated cost arbitrage without quality degradation.** Zero human intervention.

WHAT THEY BUY

- **Permission to stop paying frontier prices for trivial work —** defensible to their own risk team.

MakeMyTrip / Myra

A live production travel assistant. Not a benchmark, not a design partner pilot.

63%

LLM cost reduction

Zero

quality loss against their own rubric

One

reference that ends the “does this actually work” objection

- This is our single most valuable sales asset, and it is **already public**. It leads every first call and every cold email.
- The line that lands with a CTO: **“Your model choice is a decision you made once, nine months ago — and have paid for every day since.”**

“You're unfunded. I can't put you in the production path.”

The most common killer in enterprise DevTools, and the one we can architect away.

- **Do not argue the balance sheet. Remove the dependency.**
- Lead with **air-gapped, privately hosted deployment**: the router and the ~8B judge run inside their VPC. No traffic leaves their perimeter. **Divyam disappearing tomorrow does not take their inference down.**
- De-risk commercially: **one workload, 30-day paid pilot**, success defined as a number — ₹ saved at constant quality. Monthly terms. Written exit clause.
- Then reframe the purchase: **“You are not betting your infrastructure on us. You are renting an optimisation loop that runs on your own metal.”**

Secondary benefit: the same air-gapped posture is what unlocks BFSI. One architectural answer clears two blockers.

“We already use Portkey or LangChain.”

- Never attack the gateway. **Concede it and move the category up a layer:** “Good — keep it. That's your plumbing. I sell the decision layer above it.”

“Does your gateway automatically benchmark DeepSeek against your own production data overnight — and switch your traffic by morning, without an engineer approving it?”

- The answer is **no**, at every one of them. A human authors the routing rule. **That gap is the entire sale** — and it disqualifies fast when the account isn't ready.
- **Objection 03 — anticipate it:** “Why won't OpenAI or Bedrock just build this?” → They have a structural conflict: **no provider gets paid to route you off its own frontier model. Neutrality is the product.**

High token volume × thin unit margin × AI already in production.

No greenfield. No pilots of curiosity. Five segments, ten named accounts, one hook each.

ACCOUNT	SEGMENT	ENTRY HOOK
Swiggy / Zomato	Consumer tech	Route ~80% of trivial tier-1 support to Llama-class; reserve frontier for complex refunds. Zomato already hit cost and latency walls scaling Zia on closed models — the pain is documented and public.
Flipkart / Myntra	E-commerce	EvalMate auto-QAs millions of AI-generated catalogue descriptions against a merchandising rubric instead of human sampling. Your Flipkart tenure is the path in — used for the meeting, not the pitch.
Physics Wallah / upGrad	EdTech	Always-on AI tutors are token-hungry against low-ARPU subscriptions. Cost per active learner is a direct defence of contribution margin.
Freshworks / Darwinbox	B2B SaaS	Native copilots turn 80%-margin software into a pass-through of someone else's inference bill. Highest expansion ceiling on this list — they resell our savings to their own customers.
Cred / Groww	FinTech	Air-gapped deployment for local evaluation of sensitive financial document parsing. Sovereignty opens the door; cost closes it — strictly in that order.

Sequencing: E-commerce and B2B SaaS first (warm path, cleanest margin math) → Consumer tech second (documented pain, hardest procurement) → EdTech and FinTech run in parallel as long-cycle builds.

Two engines, deliberately unequal.

Outbound buys us revenue this quarter. Inbound buys us leverage next quarter. Confusing the two is how bootstrapped startups die.

ENGINE 1 · OUTBOUND — 80% OF MY TIME

- **Named-account, founder-assisted, technically led.** 50 accounts, ~200 contacts, hand-verified.
- Produces **all three logos** in the first 90 days.
- Cost: my time plus ~₹40k/month of tooling. **No media spend.**
- Fails if: I template it. Personalisation *is* the channel.

ENGINE 2 · INBOUND — 20% OF MY TIME

- **Technical credibility compounding into demand.** Benchmarks, teardowns, a free audit.
- Produces **near-zero pipeline in Q1** — and half of it by Q3. I will not pretend otherwise.
- Cost: content time. Distribution is borrowed, not bought.
- Fails if: we write marketing copy. Engineers forward data; they delete decks.

The honest trade: if I spend Q1 building inbound, we have beautiful content and no revenue. Outbound first is a sequencing decision, not a preference.

Five touches, twelve days, one falsifiable claim each.

THE SEQUENCE

- **D1 • Email** — one specific observation about *their* AI feature + the cost thesis. No calendar link.
- **D3 • LinkedIn** — connect with the champion. No pitch attached.
- **D5 • Email** — the MakeMyTrip 63% teardown as an artifact.
- **D8 • Email** — the disqualifying question, sent cold.
- **D12 • Break-up** — cost calculator link, permission to close the file.

MESSAGE BY PERSONA

- **VP Platform Eng** — “Here is your monthly inference bill, modelled.” *Economic pain.*
- **AI Infra Lead** — “Here is how we prove quality holds before traffic moves.” *Technical risk.*
- **CTO** — “Here is the margin story for your board.” *Strategic cover.*
- Never send the same artifact to two personas in one account.

RULES OF ENGAGEMENT

- **Tier 1 (15 accounts):** fully hand-written, plus a modelled cost estimate before first reply.
- **Tier 2 (35 accounts):** segment-templated, still manually verified.
- **Disqualify fast:** no AI in production, or under ~\$10k/month inference, or no exec sponsor by call two.
- **Volume:** 50 accounts and ~200 contacts worked per month, ~1,000 touches.

Borrow distribution. Never buy it.


Five plays, ranked by cost-to-credibility ratio. All of them are free except my time.

PLAY	MECHANIC	WHY IT CONVERTS FOR US SPECIFICALLY
Free Inference Cost Audit	30-minute teardown of their token spend, returned as a modelled savings estimate	Our highest-intent lead magnet. It is the pilot's first hour, given away. Whoever accepts is already qualified.
Quarterly benchmark reports	“DeepSeek V4 vs GPT-4o on Indian e-commerce support tickets” — real traffic shapes, published	Benchmarks are distribution in DevTools. Engineers forward them internally, which puts us in the account without a rep.
Open-source cost calculator	Small OSS tool: paste your traffic mix, see routing savings. GitHub-native	Buys developer trust that no ad can. Every star is a warm account we can name.
Comparison & intent SEO	“Portkey alternative”, “LLM cost reduction”, “LLM router” — honest, technical pages	Captures competitor-intent traffic at zero CAC , which is the only kind of paid-keyword fight we can win.
Founder-led community	Two technical posts a week, Indian ML/dev communities, Bengaluru meetups, HN and r/LocalLLaMA	The <i>model inertia</i> thesis is inherently shareable. Sandeep's voice carries further than any brand account we could build.

Instrumentation from day one: single CRM, source tagged on every lead, and I report inbound-sourced pipeline separately — so we can kill any play that isn't earning its time.

The numbers behind the 90-day commitment.

Monthly outbound engine, stated with its assumptions so you can challenge them.



WHY THESE RATES ARE DEFENSIBLE

- 10% reply assumes **hand-written, technically specific** outbound — not 2% spray-and-pray.
- 50% reply-to-meeting holds when the reply is to a **disqualifying question**, not a pitch.
- 25% pilot-to-close is conservative **because the pilot produces a number**, not an impression.

WHAT I NEED FROM YOU TO PRICE THIS

- I have **deliberately not modelled ACV** — pricing model is your call, and it changes cycle length, discounting and whether we sell to the CFO at all.
- My instinct: **a share of verified savings, floored by a platform fee**. It makes the invoice self-justifying.
- That is Question 01 on the next page.

Map, then hunt, then systematise.

DAYS 1 – 30 · BUILD & LAUNCH

Earn the right to be technical

- Map three personas: **VP Platform Engineering** (economic), **AI Infra Lead** (champion), **CTO** (air cover).
- 50 accounts → ~200 verified contacts. No purchased lists.
- Stack live: warmed inboxes, Sales Navigator, one sequencer, one CRM.
- Launch lead line: **“Cut your OpenAI inference bill by 40% this week without breaking your application.”**
- Ship inbound play 1: the free cost audit.

DAYS 31 – 60 · CONVERT

Turn interest into invoices

- Run pilots: **one workload, success criteria signed before kickoff.**
- Measure in ₹ saved at constant quality — their finance team's units, not ours.
- First benchmark report published; comparison pages live.
- Tighten ICP on evidence: double down on whichever segment replies, cut the rest.

DAYS 61 – 90 · SYSTEMATISE

Make it repeatable without me

- **3 closed logos** and one public case study cleared for use.
- POC runbook written so the next AE executes it unaided.
- Pricing validated against three real negotiations, not a spreadsheet.
- Inbound reporting its first sourced pipeline.

Six things I'd want answered in week one.

Each one changes how I sell, not just what I know.

01 · PRICING

Share of verified savings, flat platform fee, or per-token?

It sets ACV, cycle length, and whether the CFO is in the room on call one.

03 · PRODUCT EDGE

Does the router operate per-step inside agent workflows, or only per-prompt?

Agentic traffic is where 2026 token growth is. It decides whether I lead with it or avoid it.

05 · MOAT

When Portkey ships evidence-based routing, what still can't they copy?

My answer today is the reward-model IP and the customer-data flywheel. I want yours.

02 · CAPACITY

How many concurrent pilots can engineering genuinely support?

That number caps my pipeline. I'd rather pace outbound than burn a first impression.

04 · TRUST

Do we have a security questionnaire pack, SOC 2 or ISO path?

The first BFSI deal will ask in week two. Build, buy, or route around it with air-gap?

06 · DEFINITION OF WIN

At Day 90, do you want logos, revenue, or a repeatable motion?

All three is the wrong answer at this stage. Your priority sets mine.

Three things from you. All of them cheap.

- **Two warm intros from your Flipkart network in week one.** Founder-led warmth outperforms five hundred cold emails, and it compresses the first cycle by a month.
- **Engineering time-boxed to three days per pilot.** POC velocity is our only structural advantage over better-funded competitors — it is worth protecting as policy.
- **Permission to publish the cost-arbitrage numbers.** In DevTools, benchmarks are distribution — and they cost us nothing but rigour.

We don't need Portkey's budget.

We need **their customers' next invoice.**