

FOUNDING ACCOUNT EXECUTIVE · INDIA

Go-To-Market Blueprint for Divyam.AI

90-day execution strategy — turning closed-loop routing into signed enterprise revenue, without a marketing budget.

PREPARED FOR

Sandeep Kohli, CEO

PREPARED BY

Saurabh Navale

DATE

22 September 2026

Inference cost scales with usage, not headcount.

8–20×

Frontier vs open-weight price gap on the same workload

\$2.50 → \$0.30

GPT-4o vs DeepSeek-flash per 1M input tokens

~300×

Token prices fell since 2023 — enterprise bills still rose

- Nobody switches to the cheap model, because **nobody can prove quality holds**. The blocker is evidence, not price. **Model inertia**.
- **US**: heavily funded, converging on broad AI observability. Buyer is the engineer debugging quality.
- **India**: hyper cost-conscious, ROI inside a quarter. DPDP and RBI localisation gate every BFSI deal. Buyer is **the exec defending margin**.
- So: **air-gapped deployment isn't a feature — it's the entry ticket** to the top 50 Indian accounts.

Six credible players. One structural gap.

PLAYER	FUNDING & SCALE	WHERE THEY LOSE
Braintrust	Largest 2026 eval round	Scores tell you what broke. A human still acts on them.
LangSmith	LangChain-native	Locked to LangChain stacks. Observability-first, not cost-first.
TruEra / Arize	Acquired · enterprise	ML-era monitoring. Not built to move live traffic.
Portkey.ai	\$15M Series A · 500B+ tokens · 24k orgs	Routing rules are static and human-authored . Failover ≠ promotion.
Maxim AI	\$3M seed · Elevation	Evals are a destination, not a control signal.
Simplismart	\$7M A · ~\$20M Nvidia-led reported	Optimises how a model runs, never which one should.

We can't win a spend war against Portkey. **We win on architecture and on reps** — 100 surgical touches beat 100,000 impressions. Zero ad spend for 90 days.

Everyone else measures quality.

Only Divyam lets quality **automatically authorise** a cheaper model.

01 · BASELINE

EvalMate

Turns **your** examples into a rubric — not a public benchmark. Distils to an **~8B judge running on your infrastructure.**



02 · ENFORCE

Model Router

Routes every prompt across **100+ models** against your cost ceiling and quality floor. Trivial goes cheap, hard escalates.



03 · CLOSE THE LOOP

Route → Eval → Optimise

New model lands Tuesday, benchmarked on your traffic overnight, **promoted by morning — only if quality holds.**

THE PITCH

- **Automated cost arbitrage without quality degradation. Zero human intervention.**

THE PROOF

- **MakeMyTrip cut LLM cost 63% with zero quality loss** on Myra. Already public.

Two blockers kill this deal. Both have architectural answers.

01 · “YOU'RE UNFUNDED. NOT IN MY PRODUCTION PATH.”

- Don't argue the balance sheet. Remove the dependency.
- Lead with **air-gapped deployment** — router and judge inside their VPC. **If Divyam disappears, their inference doesn't.**
- One workload, **30-day paid pilot**, success = ₹ saved at constant quality. Monthly terms.
- Same answer unlocks BFSI. **One architecture, two blockers cleared.**

02 · “WE ALREADY USE PORTKEY OR LANGCHAIN.”

- Never attack the gateway. **“Keep it — that's plumbing. I sell the decision layer above it.”**
- The answer is **no** everywhere: a human authors the routing rule. **That gap is the sale.**
- **Pre-empt “won't OpenAI build this?”** — no provider gets paid to route you off its own model. **Neutrality is the product.**

The disqualifying question: “Does your gateway benchmark DeepSeek against your production data overnight and switch traffic by morning — without an engineer approving it?”

High token volume × thin margin × AI already in production.

ACCOUNT	SEGMENT	HOOK
Swiggy / Zomato	Consumer tech	Route 80% of tier-1 support to Llama-class, keep frontier for complex refunds. Zomato already hit cost walls scaling Zia.
Flipkart / Mynta	E-commerce	EvalMate auto-QAs millions of AI catalogue descriptions — machine-graded, not sampled.
Physics Wallah / upGrad	EdTech	Cut token cost per AI tutor session to protect student subscription margin.
Freshworks / Darwinbox	B2B SaaS	Embed Divyam so native copilots don't destroy software gross margin. Highest expansion ceiling.
Cred / Groww	FinTech	Air-gapped deployment for local evaluation of sensitive financial document parsing.

Sequencing: E-commerce and SaaS first (warm path, cleanest margin math) → Consumer tech second → EdTech and FinTech as parallel long-cycle builds.

Two engines, deliberately unequal.

OUTBOUND · 80% OF MY TIME · DELIVERS Q1 REVENUE

- **50 accounts, ~200 contacts**, hand-verified. Personas: VP Platform Eng, AI Infra Lead, CTO.
- **5 touches / 12 days**. Every message carries one falsifiable claim about *their* stack.
- Lead line: “**Cut your OpenAI inference bill by 40% this week without breaking your application.**”
- **Disqualify fast**: no AI in prod, under ~\$10k/mo inference, or no exec sponsor by call two.

INBOUND · 20% OF MY TIME · COMPOUNDS FROM Q2

- **Free inference cost audit** — 30-min teardown returned as a modelled savings estimate. Highest-intent magnet.
- **Quarterly benchmark reports**. In DevTools, benchmarks are distribution — engineers forward them internally.
- **OSS cost calculator** + comparison SEO (“Portkey alternative”). Zero CAC.
- **Founder-led posts**. Your voice travels further than any brand account we could build.

200

contacts / month

→

~20

replies @ 10%

→

~10

technical discos @ 50%

→

~4

paid pilots @ 40%

→

1

closed logo @ 25%

90-day
= 3 logos

Six answers that change how I sell.

01 · PRICING

Share of verified savings, platform fee, or per-token? It sets ACV and who's in the room.

03 · PRODUCT

Does the router work per-step inside agent workflows, or only per-prompt?

05 · MOAT

When Portkey ships evidence-based routing, what still can't they copy?

WHAT I NEED · 01

Two warm Flipkart-network intros in week one.

WHAT I NEED · 02

Engineering time-boxed to three days per pilot.

WHAT I NEED · 03

Permission to publish the cost-arbitrage benchmarks.

02 · CAPACITY

How many concurrent pilots can engineering support? That caps my pipeline, not my activity.

04 · TRUST

Do we have a security questionnaire pack or SOC 2 path? BFSI asks in week two.

06 · WIN

At Day 90 — logos, revenue, or a repeatable motion? Your priority sets mine.