

FOUNDING ACCOUNT EXECUTIVE · INDIA

Go-To-Market Blueprint: Divyam.AI

90-Day Execution Strategy for Founding AE

PREPARED FOR

Sandeep Kohli, CEO

PRESENTER

[My Name]

DATE

22 September 2026

Five lines. Then the evidence.

- Divyam is not an observability vendor. **Dashboards are a cost centre; Divyam is a P&L line item.**
- The wedge is the **closed loop**: EvalMate sets the quality bar, Model Router enforces it, traffic moves with no human in the loop.
- Proof already exists: **MakeMyTrip cut LLM costs 63% with zero quality loss** on Myra. That is the entire first call.
- We cannot outspend **Portkey (\$15M Series A)**. We out-architect them and out-work them on targeted outbound.
- 90-day target: **40 qualified technical discos** → **12 paid single-workload pilots** → **3 logos closed**.

Inference cost scales with usage, not headcount.

Every successful AI feature makes gross margin worse. The arbitrage is sitting unclaimed.

8–20x

Price spread between frontier and open-weight models on the same workload

~300x

Token price decline since 2023 — and enterprise bills still went **up**

\$2.50

GPT-4o per 1M input tokens vs **\$0.30** DeepSeek-flash, **\$0.10** Llama 4 Scout

- Nobody moves traffic to the cheap model because **nobody can prove quality holds**.
- That fear — not the price list — is the real product gap. Divyam already names it: **model inertia**. Most commercially useful phrase we own.

Two markets. Two entirely different buyers.

UNITED STATES

- Heavily funded, converging on **broad AI observability** — traces, spans, scores, dashboards.
- Braintrust raised one of 2026's largest eval rounds. LangSmith owns the LangChain install base.
- Their buyer is the **AI engineer debugging quality**.

INDIA

- **Hyper cost-conscious.** ROI in a quarter, not a platform vision in a year.
- **Sovereignty is contractual, not aspirational.** DPDP + RBI payment-data localisation gate every BFSI deal.
- DPDP programmes alone run **₹15–60 lakh** at large enterprises — we compete against that budget.
- Our buyer is the **exec defending margin**.

→ Consequence: **air-gapped deployment is not a feature. It is the entry ticket** to the top 50 accounts in this market.

Where each of them loses to us.

PLAYER	FUNDING / SCALE	WHERE THEY LOSE
Braintrust	Largest 2026 eval round	Scores tell you what broke. A human still fixes it.
LangSmith	LangChain-native, enterprise	Locked to LangChain-shaped stacks. Observability-first, not cost-first.
TruEra / Arize	Acquired / enterprise-scale	ML-era posture. Not built to move live GenAI traffic.
Portkey.ai	\$15M Series A · 500B+ tokens · 125M daily req · 24k orgs	Routing rules are static and human-authored . Fallbacks ≠ evidence-based promotion.
Maxim AI	\$3M seed (Elevation)	Evals are a destination , not a control signal. No traffic control plane.
Simplismart	\$7M Series A · ~\$20M Nvidia-led reported	Optimises how a model runs. Never asks which model should run.

Stated plainly, so we don't pretend otherwise.

- We will lose **every** paid-search, conference-booth and content-volume fight. Portkey has 15 million reasons to win those.
- We win on **architecture** and on **reps**. 100 surgical, technically-literate touches beat 100,000 impressions.
- Our unfair advantages are exactly three: **the closed loop**, **air-gapped deployment**, and **a POC shipped the same week it is scoped**.
- Corollary — and I mean this as a commitment: **zero top-of-funnel spend in the first 90 days**. Every rupee goes to data, tooling and POC engineering.

Everyone else measures quality.

Only Divyam lets quality **automatically authorise** a cheaper model.

01 · BASELINE

EvalMate

Agentic workflow turns **your** examples into a structured rubric — not MMLU. The aligned judge distills to an **~8B reward model running on your own infrastructure.**



02 · ENFORCE

Model Router

Lightweight classifier routes every prompt across **100+ models** against your cost ceiling and quality floor. Trivial drops cheap; hard escalates to frontier.



03 · CLOSE THE LOOP

Route → Eval → Optimize

New model lands Tuesday. Benchmarked against your traffic overnight. **Promoted by morning — only if the rubric holds.**

Automated cost arbitrage without quality degradation. Zero human intervention.

MakeMyTrip / Myra

63%

LLM cost reduction on a live production travel assistant

0

Quality loss against their own rubric

1

Reference logo that ends the "does this actually work" objection

- This is our single most valuable sales asset and it is **already public**. It leads every first call.
- The framing that lands with a CTO: **"Your model choice is a decision you made once, nine months ago — and have been paying for every day since."**

“You’re unfunded. I can’t put you in the production path.”

- Do not argue the balance sheet. Remove the dependency.
- Lead with **air-gapped / privately hosted deployment**: router and the ~8B judge run inside their VPC. No traffic leaves. **Divyam going away tomorrow does not take their inference down.**
- De-risk commercially: **one workload, 30-day paid pilot**, success = a number (₹ saved at constant quality), monthly terms, exit clause.
- Reframe: **“You are not betting your infra on us. You are renting an optimisation loop that runs on your own metal.”**

“We already use Portkey / LangChain.”

- Never attack the gateway. **Concede it and reframe the category:** “Good — keep it. That’s your plumbing. I’m selling the decision layer above it.”

“Does your current gateway automatically benchmark DeepSeek against your own production data overnight — and switch your traffic by morning, without an engineer approving it?”

- The answer is **no**. Every incumbent needs a human to author the routing rule. **That gap is the sale.**
- **BLOCKER 03 — ANTICIPATE IT** “Why won’t OpenAI or Bedrock just build this?” → Providers have a structural conflict: **they don’t get paid to route you off their frontier model. Neutrality is the product.**

High token volume × thin unit margin × AI already in production.

No greenfield. No pilots of curiosity.

ACCOUNT	SEGMENT	THE HOOK
Swiggy / Zomato	Consumer tech	Route ~80% of trivial tier-1 support to Llama-class; reserve GPT-4-class for complex refunds. Zomato already hit cost + latency walls scaling Zia on closed models — documented, live pain.
Flipkart / Myntra	E-commerce	EvalMate auto-QAs millions of AI-generated catalogue descriptions against a merchandising rubric. Sandeep's Flipkart tenure is the warm path in — use it for the meeting, not the pitch.
Physics Wallah / upGrad	EdTech	Always-on AI tutors are token-hungry against low-ARPU subscriptions. Cost per learner is a direct defence of contribution margin.
Freshworks / Darwinbox	B2B SaaS	Native copilots turn 80%-margin SaaS into a pass-through of someone else's inference bill. Strongest expansion account in the list.
Cred / Groww	FinTech	Air-gapped deployment for local evaluation of sensitive financial document parsing. Sovereignty opens the door; cost closes it — in that order.

Sequencing: Flipkart + Freshworks first (warm intro, clearest margin math) → Swiggy/Zomato second (documented pain, hardest procurement) — EdTech + FinTech as parallel long-cycle builds.

SOURCES: TOGETHER AI (ZOMATO ZIA), DATABRICKS (SWIGGY)

Map. Build. Hit the phones.

PERSONAS

- **VP / Head of Platform Eng** — owns the bill, feels it monthly. *Economic buyer.*
- **AI Infra / ML Platform Lead** — owns router and risk. *Technical champion; wins or kills it.*
- **CTO** — owns the margin story to the board. *Air cover and signature.*

BUILD · WEEK 1–2

- **50 accounts** → ~200 named **contacts**, hand-verified. No purchased lists.
- Cheap stack: inbox warm-up, Sales Nav, one sequencer, CRM. **Zero paid media.**
- Weaponise existing assets: MakeMyTrip 63%, the *model inertia* post. **Engineers forward technical content; they delete decks.**

OUTBOUND · WEEK 2–4

- Lead line: **“Cut your OpenAI inference bill by 40% this week without breaking your application.”**
- Every message carries **one falsifiable technical observation** about their stack. No “quick chat.”
- 5 touches / 12 days. **40 accounts and 200 contacts worked per week.**

What you should hold me to.

15

Technical discovery calls booked
by Day 30

6

EvalMate rubric workshops run
by Day 30

3

Paid single-workload pilots scoped
by Day 30

5

ICP messaging variants tested
by Day 30

- **D31–60:** convert pilots. One workload, measured ₹ saved at constant quality, success criteria signed **before** kickoff.
- **D61–90:** **3 closed logos**, one public case study, and a POC runbook the next AE can execute without me.

Three asks. All of them cheap.

- **Two warm Flipkart-network intros in week 1.** Founder-led warmth beats 500 cold emails.
- **Engineering time-boxed to 3 days per pilot.** POC velocity is our only structural advantage over funded competitors.
- **Permission to publish the cost-arbitrage numbers.** In DevTools, benchmarks are distribution — and they are free.

We don't need Portkey's budget.
We need **their customers' next invoice.**