

The Navier–Stokes Point

What I think you were getting at, what it says about the product, and what it says about how Divyam gets recognised.

Saurabh Navale · 25 September 2026 · For Sandeep Kohli

A caveat before anything else. This is my reconstruction of a point you made in passing, so I may have taken it somewhere you did not intend. I have written it out because it was the most interesting thing said in the conversation, and because getting it wrong in writing is more useful to both of us than nodding along in the room. Correct me where I have misread it.

1. Two readings, and they point the same way

I took the reference two ways. Both hold, and the second is the one with commercial consequences.

FIGURE 1 — THE TWO READINGS

Reading A · The classical problem

A system whose local rules are known exactly, but which has no closed-form solution. You cannot derive the answer in advance. You can only step forward numerically, measure, and correct. Turbulence makes long-horizon prediction structurally impossible — not merely difficult.

→

Reading B · The September result

On 8 September 2026 OpenAI published a proposed resolution of the existence and smoothness problem, produced by a large internal multi-agent system and formalised in Lean. Clay still lists the problem as open; OpenAI did not claim the prize.

2. Reading A — Divyam is a numerical solver, not a rules engine

If the analogy is the classical problem, then it is a precise description of what model routing actually is.

- **There is no closed form for "which model should answer this prompt."** You cannot derive it from model cards, benchmark tables or parameter counts. The only honest method is to run it against real traffic and measure the result.
- **The field is turbulent, in the technical sense.** New weights land weekly, prices move, quality drifts, traffic mix shifts with the product. Any routing rule written by hand is a snapshot of a flow that has already moved on.
- **So a static gateway is the wrong shape of answer.** It is an analytical solution offered for a problem that does not admit one. It will be right on the day it is written and progressively wrong every day after.

The reframe I would take into a sales conversation. Competitors ship a rules engine and ask a human to keep it current. Divyam ships a solver: it steps the system forward, measures against a quality bar defined on the customer's own data, and corrects. **The closed loop is not a feature on top of routing. It is the only method that fits the shape of the problem.**

It also sharpens the phrase we already use. *Model inertia* is not laziness on the customer's part. It is the rational response to being handed a system they cannot solve analytically and given no instrument to measure it with.

3. Reading B — how a deep-tech company actually gets recognised

DIVYAM.AI · FOLLOW-UP NOTE

SAURABH NAVALE · 25 SEP 2026

This is the part I keep returning to. In the first week of September, a company bought an enormous amount of market attention with a single result. Not with a campaign, a launch event or a category definition. The mechanics of that are worth copying precisely, because they are available to a bootstrapped company and a marketing budget is not.

WHAT MADE IT LAND	WHY IT TRANSFERS TO DIVYAM
The claim was hard and specific	Not "we are good at maths." A named problem, stated by someone else, that the field already agreed was difficult. Our equivalent: a named workload, a named model pair, a named cost delta.
It was independently checkable	The Lean formalisation is the whole trick: you do not have to trust the claimant. Our equivalent is a published method and a reproducible harness — not a case-study PDF.
They declined the prize	Refusing the obvious reward signalled that the result, not the trophy, was the point. Restraint reads as confidence. Worth remembering when we are tempted to over-claim a number.
The result did the distribution	Nobody needed to be persuaded to pass it on. In developer infrastructure, a verifiable benchmark travels further than any campaign we could fund.

The uncomfortable half of the lesson. That result also came from an organisation with effectively unlimited compute. We cannot copy the scale. We can copy the **form**: one hard, specific, independently verifiable claim, published in a way that an engineer can rerun. That is the only recognition mechanism that does not have a budget attached to it.

4. What I would actually do with this

Four moves, in the order I would run them. All are inside the 90-day plan I sent; this is what the Navier–Stokes framing changes about how they are positioned.

#	MOVE	THE POINT OF IT
1	A public, reproducible index	Quarterly: cost-versus-quality frontier across named models on named task families — Indian e-commerce support tickets, catalogue generation, document parsing. Published with the method, so anyone can rerun it. This is our Lean file.
2	An open measurement harness	The thing that makes claim one checkable. It also quietly makes Divyam's definition of quality the one the market argues in. Whoever owns the measurement owns the category.
3	Customer results stated as claims, not testimonials	"63% at constant quality on this workload, measured this way, over this period." A number someone could attack is worth more than three that nobody can check.
4	Name the category as a method	Not "AI gateway" — that is Portkey's word and their shape. Something closer to closed-loop inference control : the claim that this class of problem is solved by continuous measurement and correction, and by nothing else.

5. The corollary I would not skip

DIVYAM.AI · FOLLOW-UP NOTE

SAURABH NAVALE · 25 SEP 2026

Note *how* the September result was produced: a large multi-agent system, run long. That is the shape of serious AI workloads now, and it has a direct commercial consequence for us.

- Agentic workloads multiply token volume by every step in the chain. **A single user request becomes dozens of model calls.**
- Most of those steps are routine — retrieval, formatting, checking, summarising — and are wildly overserved by a frontier model.
- Which means the cost problem we sell into is **getting worse on a curve, not a line**, and the buyer's urgency rises with it.

This is why I asked whether the router operates per step inside agent workflows or only per request. If it is per step, the Navier–Stokes story is not an analogy we borrow — it is a description of the exact workload class where Divyam's value compounds fastest. I would lead with that.

6. Where I might be wrong

- You may have meant something narrower — flow control, or turbulence as a pure metaphor for market churn — and I have built more on it than it was meant to carry.
- The index in move one costs real engineering time. If pilot capacity is the binding constraint, publishing a benchmark competes with servicing customers, and customers win.
- Naming a category is usually a thing a market does to you rather than a thing you do to a market. I would earn it with moves one to three before claiming it in move four.

The question I would put back to you. If Divyam could publish exactly one result in the next ninety days — one hard, specific, independently verifiable number — which workload should it be about? I would build the first index around that answer, and I would use it as the opening artifact for every account on the target list.

Thank you for the conversation. It was the first interview I have had where the most interesting sentence was one I had to go away and think about. Whatever you decide, the thinking was worth doing.

Saurabh Navale

Sources for the September result: OpenAI, "On the Navier–Stokes Millennium Prize Problem" (8 Sep 2026); Quanta Magazine, 8 Sep 2026; Nature news, Sep 2026; Clay Mathematics Institute, Millennium Prize Problems — the problem remains formally listed as open.