

What Divyam Sells, and How I Would Sell It

Following our conversation — the problem in plain terms, the point you raised about Navier–Stokes, and the outreach plan.

Saurabh Navale · 25 September 2026

What this document is. Three pages. Page one is the problem we sell into. Page two is what Divyam does about it. Page three is how I would go and find the buyers. No technical detail — this is written the way I would explain it to a customer.

01 The problem, in five lines

- Companies run their AI features on the most expensive models available.
- Most of what they send is routine work — lookups, sorting, formatting, summarising.
- A cheaper model would handle that routine work just as well.
- **Nobody switches, because nobody can prove the quality will hold.**
- So the bill keeps growing, and their data keeps leaving the building to get processed.

02 Why "sending it outside" stopped being acceptable

- Research, financial documents and customer records leave the company to be processed.
- This year the AI labs themselves disclosed incidents where their own models reached systems they were never meant to touch.
- The risk stopped being theoretical, and buyers noticed.
- **For banks and universities it is no longer a preference. It is a rule they have to follow.**

03 The Navier–Stokes point you raised

- Navier–Stokes is the famous maths problem that **cannot be solved on paper** — you can only measure the system as it moves and keep correcting.
- **Choosing which model should answer which request is the same kind of problem.** There is no formula. Model rankings and price lists will not tell you.
- Which is why competitors get it wrong: they hand the customer a set of fixed rules, written by a person, on a day. New models arrive weekly, so those rules are out of date almost immediately.

The sentence I would use with a customer: “Everyone else gives you rules that are stale the day they are written. Divyam measures your actual traffic and corrects itself — continuously.”

HOW IT WORKS — THREE STEPS

1. Learn their standard

Divyam studies the customer's own examples to learn what a good answer looks like for them — not a generic score.

→

2. Send work to the cheapest model that passes

Every single request is matched to the cheapest model that still meets that standard. Hard requests still go to the expensive model.

→

3. Keep checking, keep correcting

New model arrives, traffic changes, quality drifts — Divyam re-tests and adjusts on its own.

WHERE THE SAVING COMES FROM — THE WORK MIX, ESTIMATED

Routine work — roughly three quarters

Genuinely hard — the rest

Today **all** of this runs on the expensive model. Divyam moves the left-hand block to cheaper models and leaves the right-hand block alone.
Estimated outcome: roughly half the cost, with no change to the product and no drop in quality.

- **It runs inside the customer's own environment** — nothing leaves their building.
- **Nothing moves unless quality holds.** The standard is the gate, not a hope.
- **No engineering work for the customer** — the product stays exactly as it is.
- **Already proven.** MakeMyTrip cut their AI cost by more than half on a live product with no quality loss.

05 The GPU partnership — a second route to market

- Tie-ups with Indian GPU providers such as **Yotta and E2E**, including a place on their marketplace.
- They sell the hardware. Customers often do not know what is genuinely worth running on it.
- **Divyam makes that hardware useful** by proving the cheaper open models are good enough to rely on.
- Their sales teams reach buyers who have already spent the money — that is distribution we do not pay for.
- Especially useful for universities, who often buy their compute from exactly these providers.

A note on this split. You raised universities as an example, not a decided plan. Below is the weight I would give it — a recommendation, not an assumption. The same plan works at any ratio.

OUTREACH EFFORT — PROPOSED

UNIVERSITIES — 40%

CORPORATE — 60%

2 days a week · ~20 institutions a month

3 days a week · ~30 companies a month

WHY UNIVERSITIES GET 40%

- Large numbers of users in a single institution.
- Keeping data in-house is written policy, not preference.
- Tight budgets make “half the cost” an easy conversation.
- **Nobody else is calling on them.**
- Students become users who carry it into industry later.

WHY CORPORATE STAYS 60%

- Shorter decisions, tied to a quarter.
- This is where revenue lands inside the funding window.
- Bigger, faster deals once a single workload is proven.

07 How each side is worked

	UNIVERSITIES	CORPORATE
Who I contact	Dean of Research, CIO, Head of Computer Centre — and lab heads directly	Engineering leadership, and whoever owns the AI budget
What I open with	An estimate of what their current setup costs them per experiment	An estimate of their monthly AI bill, worked out before the call
Lead argument	Data stays in-house first , cost second	Cost first , data second
Cheapest way in	A technical talk on campus — fills a room no email would reach	Being roughly right about their bill before they tell me

08 The first 90 days

	UNIVERSITIES	CORPORATE
Month 1	Build the institution list, book two speaking slots, start conversations with individual labs	Build the account list, open cost conversations, first discovery calls
Month 2	Take one institution up to Dean or CIO level, scope a trial with a single lab	Two to three trials scoped and signed, one workload each
Month 3	One institutional trial running and producing a result we can publish	First closes — with the customer’s own team presenting the result internally

WHAT I TRACK

- **How many accounts I can state a current AI spend for** — everything follows this one.
- Trials scoped, and trials running.
- Users activated on each side.

WHAT GOOD LOOKS LIKE AT DAY 90

- About a third of the pipeline from universities.
- Most of the closed revenue from corporate.
- One published university result — that travels further than any advertising we could buy.
- At least one GPU partner sending introductions.

In one line: Divyam cuts the AI bill roughly in half without the customer’s product changing and without their data leaving the building — and my recommendation would be to give universities a real share of the week, because that is where the need is written into policy and nobody else is knocking.