

What Divyam Sells, and How I Would Sell It

Following our conversation — the problem in plain terms, the point you raised about Navier–Stokes, and the outreach plan.

Saurabh Navale · 25 September 2026

What this document is. Three pages. Page one is the problem we sell into. Page two is what Divyam does about it. Page three is how I would go and find the buyers. No technical detail — this is written the way I would explain it to a customer.

01 The problem, in five lines

- Companies run their AI features on the most expensive models available.
- Most of what they send is routine work — lookups, sorting, formatting, summarising.
- A cheaper model would handle that routine work just as well.
- **Nobody switches, because nobody can prove the quality will hold.**
- So the bill keeps growing, and their data keeps leaving the building to get processed.

02 Why "sending it outside" stopped being acceptable

- Research, financial documents and customer records leave the company to be processed.
- This year the AI labs themselves disclosed incidents where their own models reached systems they were never meant to touch.
- The risk stopped being theoretical, and buyers noticed.
- **For banks and universities it is no longer a preference. It is a rule they have to follow.**

03 The Navier–Stokes point you raised

- Navier–Stokes is the famous maths problem that **cannot be solved on paper** — you can only measure the system as it moves and keep correcting.
- **Choosing which model should answer which request is the same kind of problem.** There is no formula. Model rankings and price lists will not tell you.
- Which is why competitors get it wrong: they hand the customer a set of fixed rules, written by a person, on a day. New models arrive weekly, so those rules are out of date almost immediately.

The sentence I would use with a customer: “Everyone else gives you rules that are stale the day they are written. Divyam measures your actual traffic and corrects itself — continuously.”

04 What Divyam actually does

HOW IT WORKS — THREE STEPS

1. Learn their standard

Divyam studies the customer's own examples to learn what a good answer looks like for them — not a generic score.



2. Send work to the cheapest model that passes

Every single request is matched to the cheapest model that still meets that standard. Hard requests still go to the expensive model.



3. Keep checking, keep correcting

New model arrives, traffic changes, quality drifts — Divyam re-tests and adjusts on its own.

WHERE THE SAVING COMES FROM — THE WORK MIX, ESTIMATED

Routine work — roughly three quarters

Genuinely hard — the rest

Today **all** of this runs on the expensive model. Divyam moves the left-hand block to cheaper models and leaves the right-hand block alone.
Estimated outcome: roughly half the cost, with no change to the product and no drop in quality.

- **It runs inside the customer's own environment** — nothing leaves their building.
- **Nothing moves unless quality holds.** The standard is the gate, not a hope.
- **No engineering work for the customer** — the product stays exactly as it is.
- **Already proven.** MakeMyTrip cut their AI cost by more than half on a live product with no quality loss.

05 The GPU partnership — a second route to market

- Tie-ups with Indian GPU providers such as **Yotta and E2E**, including a place on their marketplace.
- They sell the hardware. Customers often do not know what is genuinely worth running on it.
- **Divyam makes that hardware useful** by proving the cheaper open models are good enough to rely on.
- Their sales teams reach buyers who have already spent the money — that is distribution we do not pay for.
- Especially useful for universities, who often buy their compute from exactly these providers.

06 How I would split my time

OUTREACH EFFORT

UNIVERSITIES — 40%

CORPORATE — 60%

2 days a week · ~20 institutions a month

3 days a week · ~30 companies a month

WHY UNIVERSITIES GET 40%

- Large numbers of users in a single institution.
- Keeping data in-house is written policy, not preference.
- Tight budgets make “half the cost” an easy conversation.
- **Nobody else is calling on them.**
- Students become users who carry it into industry later.

WHY CORPORATE STAYS 60%

- Shorter decisions, tied to a quarter.
- This is where revenue lands inside the funding window.
- Bigger, faster deals once a single workload is proven.
- Produces the reference names universities respect.

07 How each side is worked

	UNIVERSITIES	CORPORATE
Who I contact	Dean of Research, CIO, Head of Computer Centre — and lab heads directly	Engineering leadership, and whoever owns the AI budget
What I open with	An estimate of what their current setup costs them per experiment	An estimate of their monthly AI bill, worked out before the call
Lead argument	Data stays in-house first , cost second	Cost first , data second
Cheapest way in	A technical talk on campus — fills a room no email would reach	Being roughly right about their bill before they tell me
Deal shape	Whole institution, annual, slower	One team first, then spreads

08 The first 90 days

	UNIVERSITIES	CORPORATE
Month 1	Build the institution list, book two speaking slots, start conversations with individual labs	Build the account list, open cost conversations, first discovery calls
Month 2	Take one institution up to Dean or CIO level, scope a trial with a single lab	Two to three trials scoped and signed, one workload each
Month 3	One institutional trial running and producing a result we can publish	First closes — with the customer’s own team presenting the result internally

WHAT I TRACK

- **How many accounts I can state a current AI spend for** — everything follows this one.
- Trials scoped, and trials running.
- Users activated on each side.

WHAT GOOD LOOKS LIKE AT DAY 90

- About a third of the pipeline from universities.
- Most of the closed revenue from corporate.
- One published university result — that travels further than any advertising we could buy.
- At least one GPU partner sending introductions.

In one line: Divyam cuts the AI bill roughly in half without the customer’s product changing and without their data leaving the building — and I would spend four days in ten on universities, because that is where the need is written into policy and nobody else is knocking.